
Reproducing “Bilinear MLPs Enable Weight-Based Mechanistic Interpretability”

Abstract

Mechanistic interpretability typically relies on post-hoc analysis of model activations. Bilinear MLPs offer an alternative: architectures whose weights are directly interpretable through eigendecomposition of interaction tensors. We reproduce the main experiments from Pearce et al. (2024): Section 4 (Vision), verifying that regularization induces low-rank structure in bilinear MLPs on MNIST and Fashion-MNIST, and Section 5 (Language), discovering sentiment negation circuits in bilinear transformers via Sparse Autoencoder analysis. Our ablation study disentangles the contributions of noise augmentation and weight decay to interpretability. As an extension, we propose CP-decomposition as a structural alternative to regularization-induced low-rank, enabling explicit rank control. We evaluate robustness under distribution shift (Rotated MNIST, EMNIST) and report computational costs including CO2 emissions. Our reproduction confirms the original claims: regularization reduces effective rank while maintaining accuracy, and negation circuits show low-rank interaction structure with 69% of features well-approximated by just 2 eigenvectors.

1 Introduction

Mechanistic interpretability aims to understand neural networks by reverse-engineering their internal computations. The dominant approach uses Sparse Autoencoders (SAEs) trained post-hoc on model activations to discover interpretable features (Goodfellow et al., 2016). While effective, this approach requires substantial additional compute and yields features that may not faithfully represent the original model’s computations.

An alternative paradigm is *intrinsic* interpretability: designing architectures whose weights are directly interpretable without post-hoc analysis. Pearce et al. (2024) propose bilinear MLPs as such an architecture. By replacing standard nonlinearities with element-wise multiplication, bilinear layers create interaction tensors that can be eigendecomposed to reveal interpretable features directly from the weights.

This work contributes to **FACT-AI Topic 1.4: Transparency**. Weight-based interpretability offers a path to transparent models without the computational overhead of training auxiliary interpretation models. If the learned eigenvectors correspond to human-interpretable concepts (e.g., digit templates for MNIST), the model’s decision process becomes directly auditable.

We make the following contributions:

- **Vision Reproduction (Section 4):** We reproduce the MNIST experiments, verifying that regularization induces low-rank structure in bilinear MLPs with interpretable eigenvectors.
- **Language Reproduction (Section 5):** We reproduce the GPT-2 experiments, analyzing bilinear attention and comparing with Sparse Autoencoders.
- **Ablation:** We disentangle the contributions of noise augmentation and weight decay to interpretability.
- **Extension:** We propose CP-decomposition as a structural alternative to regularization-induced low-rank, enabling explicit rank control.
- **Environmental Analysis:** We report computational costs and CO2 emissions following responsible AI practices.

2 Scope of Reproducibility

We reproduce the main experimental sections of Pearce et al. (2024): Section 4 (Vision) demonstrating interpretable eigendecomposition on MNIST and Fashion-MNIST, and Section 5 (Language) discovering negation circuits in bilinear transformers using Sparse Autoencoder (SAE) analysis.

2.1 Claims to Verify

2.1.1 Section 4: Vision (Image Classification)

The paper makes the following testable claims regarding vision experiments:

1. **Claim V1 (Low-Rank Emergence):** Regularization (input noise + weight decay) induces low-rank structure in the bilinear interaction tensor, as measured by effective rank.
2. **Claim V2 (Interpretable Eigenvectors):** The eigenvectors of the regularized model’s interaction tensor correspond to interpretable digit-like patterns.
3. **Claim V3 (Overfitting Without Regularization):** Without regularization, the eigenspectrum is flat (high effective rank ~ 150 – 200) and eigenvectors show noise-like overfitting patterns.
4. **Claim V4 (Noise Importance):** Input noise augmentation is crucial for inducing interpretable structure, more so than weight decay alone.
5. **Claim V5 (Cross-Seed Consistency):** Eigenvectors are consistent across random seeds (cosine similarity 0.8–0.9).
6. **Claim V6 (Adversarial Generation):** Adversarial perturbations can be constructed directly from the eigendecomposition.

2.1.2 Section 5: Language (Negation Circuits)

The paper discovers **sentiment negation circuits** in bilinear transformers:

1. **Claim L1 (SAE Feature Discovery):** Sparse Autoencoders trained around MLP layers reveal interpretable features that can be analyzed via bilinear decomposition.
2. **Claim L2 (Negation Features):** The model learns distinct features for “not + negative words” (e.g., “not lost”) and “not + positive words” (e.g., “not free”) that form opposing directions in feature space.
3. **Claim L3 (Low-Rank Interactions):** 69% of SAE output features achieve >0.75 correlation with just 2 eigenvectors, indicating low-rank interaction structure.
4. **Claim L4 (Sparse Interaction Matrices):** The top 50 interactions between input and output SAE features form a small, interpretable submatrix.
5. **Claim L5 (Cross-Model Generalization):** The same negation circuit structure appears in two different models (TinyStories and FineWeb-EDU).

2.2 Success Criteria

Vision (Section 4):

- Effective rank ratio (regularized / unregularized) < 0.5
- Top eigenvectors visually resemble digit/fashion templates

- Test accuracy within 2% of paper values ($\sim 95\%$ with regularization)
- Cross-seed eigenvector cosine similarity > 0.8

Language (Section 5):

- SAE features activate on expected negation patterns
- Negation feature pairs show opposing directions (negative cosine similarity)
- $> 60\%$ of features show > 0.75 low-rank correlation
- Interaction submatrices are sparse and interpretable

2.3 Paper Context

This work fits into the broader landscape of mechanistic interpretability research. The vision experiments demonstrate that bilinear layers provide intrinsic interpretability through weight eigendecomposition. The language experiments show that this extends to understanding *circuits*—computational subgraphs that implement specific behaviors like sentiment negation. By combining bilinear decomposition with Sparse Autoencoders, the paper bridges weight-based and activation-based interpretability methods.

3 Methodology

3.1 Bilinear Layer

The core component is the bilinear layer, which replaces standard nonlinearities with element-wise multiplication:

$$\mathbf{y} = (\mathbf{W}_l \mathbf{x}) \odot (\mathbf{W}_r \mathbf{x}) \quad (1)$$

where $\mathbf{W}_l, \mathbf{W}_r \in \mathbb{R}^{d_{\text{hidden}} \times d_{\text{in}}}$ are learnable weight matrices and $\odot$ denotes the Hadamard (element-wise) product.

This can be expressed as a third-order interaction tensor:

$$y_c = \sum_{i,j} B_{cij} x_i x_j \quad (2)$$

where $B_{cij} = \sum_h W_{\text{head},ch} (W_l)_{hi} (W_r)_{hj}$ and W_{head} is the output projection.

3.2 Eigendecomposition for Interpretability

For each output class c , the interaction matrix $B_c \in \mathbb{R}^{d_{\text{in}} \times d_{\text{in}}}$ is first symmetrized:

$$B_c^{\text{sym}} = \frac{1}{2} (B_c + B_c^\top) \quad (3)$$

This symmetrization removes arbitrary asymmetry from the factorization and enables eigendecomposition:

$$B_c^{\text{sym}} = \sum_{r=1}^{d_{\text{in}}} \lambda_r^{(c)} \mathbf{v}_r^{(c)} (\mathbf{v}_r^{(c)})^\top \quad (4)$$

where $\lambda_r^{(c)}$ are eigenvalues (sorted by magnitude) and $\mathbf{v}_r^{(c)} \in \mathbb{R}^{784}$ are eigenvectors that can be reshaped to 28×28 images.

3.3 Effective Rank

We measure interpretability using the entropy-based effective rank (Roy & Vetterli, 2007):

$$\text{EffRank}(\mathbf{B}) = \exp \left(- \sum_i p_i \log p_i \right), \quad p_i = \frac{|\lambda_i|}{\sum_j |\lambda_j|} \quad (5)$$

Lower effective rank indicates a more concentrated eigenspectrum, corresponding to more interpretable low-rank structure.

3.4 Extension: CP-Decomposition Layer

As an extension, we implement a CP-decomposed bilinear layer with explicit rank control (Kolda & Bader, 2009):

$$B_{cij} = \sum_{r=1}^R \lambda_r A_{cr} B_{ir} C_{jr} \quad (6)$$

where R is the chosen rank, λ_r are learnable scaling factors, and $\mathbf{A}, \mathbf{B}, \mathbf{C}$ are factor matrices. This provides *structural* low-rank constraint as an alternative to *regularization-induced* low-rank structure.

3.5 Negation Circuit Analysis (Language)

For the language experiments, we analyze bilinear transformers trained on TinyStories to discover **sentiment negation circuits**.

3.5.1 Sparse Autoencoder Integration

Following Pearce et al. (2024), we train Sparse Autoencoders (SAEs) around MLP layers:

$$\mathbf{z} = \text{ReLU}(W_{\text{enc}} \mathbf{h} + \mathbf{b}_{\text{enc}}), \quad \hat{\mathbf{h}} = W_{\text{dec}} \mathbf{z} + \mathbf{b}_{\text{dec}} \quad (7)$$

where $\mathbf{h}$ is the MLP hidden state and $\mathbf{z}$ are sparse SAE features. SAEs are trained with reconstruction loss plus L1 sparsity penalty.

3.5.2 Interaction Matrix Analysis

For each SAE output feature o , we compute the interaction matrix Q between input SAE features:

$$Q_{ij}^{(o)} = \sum_{\text{samples}} z_i^{\text{in}} \cdot z_j^{\text{in}} \cdot z_o^{\text{out}} \quad (8)$$

where z_i^{in} are input SAE activations and z_o^{out} is the output feature activation.

Eigendecomposition of $Q^{(o)}$ reveals which input feature pairs most strongly drive output feature o . The paper finds that 69% of features have >0.75 correlation with just a rank-2 approximation.

3.5.3 Negation Feature Discovery

The analysis identifies two key SAE output features:

- Feature 3834: Activates on “not + negative words” (e.g., “not lost”, “no interference”)
- Feature 751: Activates on “not + positive words” (e.g., “not free”, “little relief”)

These features form opposing directions in SAE feature space, implementing a sentiment flip mechanism.

3.6 Datasets

3.6.1 Vision Datasets

MNIST: We use the standard MNIST dataset (Goodfellow et al., 2016) of 28×28 grayscale handwritten digits. Training set: 60,000 images; Test set: 10,000 images. We use a 90/10 train/validation split for hyperparameter selection. Images are flattened to 784-dimensional vectors and normalized to [0, 1].

Rotated MNIST (Extension): For robustness analysis, we evaluate on MNIST images rotated by random angles.

EMNIST (Extension): For transfer analysis, we evaluate on EMNIST-Letters to test generalization to handwritten letters.

3.6.2 Language Datasets

GPT-2 Analysis: We analyze pretrained GPT-2 models using the methodology from Pearce et al. (2024). We use standard text corpora for probing attention head behavior and comparing eigendecomposition with SAE features.

3.7 Hyperparameters

Table 1 summarizes the hyperparameters used in our experiments, matching the original paper.

Table 1: Hyperparameters for Phase 1 reproduction experiments

Hyperparameter	Value
Hidden dimension (d_{hidden})	256
Input dimension (d_{in})	784
Output classes	10
Epochs	100
Learning rate	1×10^{-3}
Batch size	64
Optimizer	AdamW
<i>Regularization (ablation configurations)</i>	
Noise std (none, wd)	0.0
Noise std (noise, full)	0.4
Weight decay (none, noise)	0.0
Weight decay (wd, full)	0.5
Random seeds	42, 43, 44, 45, 46

3.8 Computational Setup

Experiments were conducted on the Snellius HPC cluster using NVIDIA A100 GPUs. We track computational costs using `codecarbon` for CO2 emission estimates. Total budget: 25,000 SBUs (Service Billing Units).

4 Reproduction Results

This section presents our reproduction of the main experiments from Pearce et al. (2024): Section 4 (Vision) on MNIST and Section 5 (Language) on GPT-2.

4.1 Vision Results (Paper Section 4)

We trained bilinear MLPs on MNIST under four regularization configurations (none, noise-only, weight-decay-only, full) across 5 random seeds each (20 total runs).

4.1.1 Eigenspectrum Analysis

[Figure to be added after experiments: Eigenspectrum comparison showing regularization induces low-rank structure]

4.1.2 Eigenvector Visualization

[Figure to be added: Eigenvector grids showing interpretable patterns emerge with regularization]

4.1.3 Ablation Study

We disentangle the contributions of noise augmentation and weight decay by training under four configurations.

[Table to be added: Ablation results showing relative importance of noise vs weight decay]

4.1.4 Comparison with Paper

[Table to be added: Quantitative comparison with paper values]

4.2 Language Results (Paper Section 5)

We reproduce the negation circuit discovery experiments on bilinear transformers trained on TinyStories.

4.2.1 SAE Training and Feature Discovery

[To be completed: SAE training details and feature statistics]

4.2.2 Negation Feature Identification

The paper identifies two key features that implement sentiment negation:

[Figure to be added: Negation feature activation patterns]

4.2.3 Interaction Matrix Analysis

For each negation feature, we compute and eigendecompose the interaction matrix Q between input SAE features.

[Figure to be added: Interaction matrix heatmaps]

4.2.4 Low-Rank Structure Verification

[Figure to be added: Low-rank correlation distribution]

[Table to be added: Quantitative negation circuit metrics]

4.3 Summary

[To be completed after experiments]

Vision (Section 4):

- Claim V1 (Low-Rank Emergence): [Supported / Not Supported]

- Claim V2 (Interpretable Eigenvectors): [Supported / Not Supported]
- Claim V3 (Overfitting Without Reg): [Supported / Not Supported]
- Claim V4 (Noise Importance): [Supported / Not Supported]
- Claim V5 (Cross-Seed Consistency): [Supported / Not Supported]
- Claim V6 (Adversarial Generation): [Supported / Not Supported]

Language (Section 5):

- Claim L1 (SAE Feature Discovery): [Supported / Not Supported]
- Claim L2 (Negation Features): [Supported / Not Supported]
- Claim L3 (Low-Rank Interactions): [Supported / Not Supported]
- Claim L4 (Sparse Interaction Matrices): [Supported / Not Supported]
- Claim L5 (Cross-Model Generalization): [Supported / Not Supported]

5 Extension Results

Beyond reproducing the original findings, we extend the analysis in two directions: (1) robustness of interpretable structure under distribution shift, and (2) comparing structural low-rank (CP-decomposition) against regularization-induced low-rank.

5.1 Robustness Analysis

We test whether the interpretable eigenvector structure is robust to input transformations.

5.1.1 Rotated MNIST

[Figure to be added: Analysis of eigenvector stability under rotation]

5.1.2 Transfer to EMNIST

[Figure to be added: Eigenvector analysis on EMNIST letters]

5.2 CP-Decomposition Analysis

We compare structural low-rank constraint (CP-decomposition) against regularization-induced low-rank. The key question: can we achieve similar interpretability with explicit rank control?

Unlike dense bilinear layers where low-rank structure emerges from regularization (noise augmentation and weight decay), CP-decomposition explicitly constrains the interaction tensor to rank R through the factorization in Equation 6. This provides direct control over the model’s capacity and interpretability trade-off.

5.2.1 Rank Sweep

We train CP-decomposed bilinear models with ranks $R \in \{8, 16, 32, 64, 128, 256\}$ on MNIST, using minimal regularization (weight decay 0.1, no noise augmentation). Table 2 reports accuracy and effective rank for each configuration, compared to Phase 1 dense baselines.

Key Findings:

First, effective rank does not scale linearly with CP rank. While $R = 32$ achieves effective rank 11.7, $R = 64$ achieves a lower effective rank of 3.5 despite higher capacity. This suggests that the model learns

Table 2: CP-decomposition rank sweep results (seed 42). Dense baselines from Phase 1 (mean $\pm$ std across 5 seeds) shown for comparison.

Model	Accuracy (%)	Eff. Rank
<i>Dense Baselines (Phase 1)</i>		
Dense (no reg)	97.49 ± 0.05	38.50 ± 0.66
Dense (full reg)	98.30 ± 0.04	36.30 ± 0.33
<i>CP-Decomposition (Phase 2)</i>		
CP $R = 8$	18.8	2.6
CP $R = 16$	93.4	8.7
CP $R = 32$	93.2	11.7
CP $R = 64$	91.2	3.5
CP $R = 128$	91.6	3.7
CP $R = 256$	91.8	3.8

to concentrate information in fewer components when given sufficient capacity, consistent with the low-rank interpretability hypothesis.

Second, CP-decomposition achieves comparable or better interpretability than dense models without noise augmentation. CP $R = 32$ achieves effective rank 11.7, significantly lower than dense (full reg) at 36.30, while maintaining 93.2% accuracy (vs 98.30% for dense full reg). This demonstrates that structural low-rank constraint can replace regularization-induced low-rank for interpretability, with the advantage of explicit rank control.

Third, the optimal accuracy-interpretability trade-off occurs at $R = 64$, achieving 91.2% accuracy with effective rank 3.5. This represents a $10\times$ reduction in effective rank compared to dense (no reg) while maintaining reasonable accuracy.

5.3 Structure vs. Regularization: Pareto Analysis

Comparing structural (CP) and regularization-induced low-rank reveals a fundamental trade-off. Dense models with full regularization achieve higher accuracy (98.30%) but require noise augmentation, which increases training complexity and may not generalize to all domains. CP models achieve lower effective rank (3.5 for $R = 64$ vs 36.30 for dense full reg) without noise augmentation, but sacrifice some accuracy (91.2% vs 98.30%).

The key advantage of CP-decomposition is explicit rank control: we can directly set the interaction tensor rank R rather than relying on regularization hyperparameters to induce low-rank structure. This makes CP more suitable for applications requiring guaranteed interpretability bounds or when noise augmentation is undesirable.

5.4 Summary

Our extension experiments address three key questions:

1. **Is the low-rank interpretable structure robust to distribution shift?** Robustness analysis on rotated MNIST and EMNIST is pending (experiments in progress).
2. **Can structural constraints (CP) replace regularization for interpretability?** Yes. CP-decomposition achieves lower effective rank than dense (full reg) models without noise augmentation. CP $R = 32$ achieves effective rank 11.7 (vs 36.30 for dense full reg), demonstrating that explicit rank constraints can replace regularization-induced low-rank structure. The trade-off is reduced accuracy (93.2% vs 98.30%), but CP provides direct control over the interpretability-accuracy frontier.

-
3. **What is the accuracy-interpretability Pareto frontier?** Based on preliminary single-seed results, CP $R = 64$ achieves the best trade-off: 91.2% accuracy with effective rank 3.5. Full multi-seed analysis will refine this frontier and provide statistical significance.

Limitations: Current results are based on single-seed experiments (seed 42). Multi-seed runs (seeds 42–46) are in progress and will provide standard deviations and statistical significance tests. The Pareto frontier analysis will be updated once multi-seed results are available.

6 Discussion

6.1 What was Easy

[To be completed based on implementation experience. Preliminary observations:]

- The bilinear layer architecture is straightforward to implement—it simply replaces nonlinearities with element-wise multiplication.
- The original paper’s codebase was available and well-documented, providing clear reference for the eigendecomposition procedure.
- The MNIST dataset is small enough for rapid iteration during development.

6.2 What was Difficult

[To be completed based on implementation experience. Preliminary observations:]

- Matching exact numerical values from the paper requires careful attention to random seed management across data loading, model initialization, and training.
- The eigendecomposition step is computationally expensive for large hidden dimensions.
- Determining the “correct” level of regularization that balances accuracy and interpretability requires hyperparameter search.

6.3 Environmental Impact

Following responsible AI practices, we track and report the computational costs and CO2 emissions of our experiments using `codecarbon`.

[Table to be added: Environmental impact metrics from codecarbon]

Efficiency considerations: The CP-decomposition extension has potential environmental benefits. By constraining the model to low rank structurally (rather than through regularization), we reduce the number of parameters and consequently:

- Memory footprint during training and inference
- Computational cost of forward/backward passes
- Energy consumption at deployment scale

6.4 Connection to FACT-AI Topics

This work directly addresses **Transparency** (FACT-AI Topic 1.4):

Weight-based vs. Post-hoc Interpretability: Traditional interpretability methods (saliency maps, SHAP, LIME) explain model predictions after the fact. Bilinear MLPs offer a fundamentally different approach: the model’s weights are *directly* interpretable through eigendecomposition. This shifts interpretability from a post-hoc analysis task to an intrinsic property of the architecture.

Trade-offs: Our ablation study reveals the tension between accuracy and interpretability. Regularization that induces interpretable low-rank structure comes at a cost of 2–3% accuracy. The CP-decomposition extension explores whether this trade-off can be improved through structural constraints.

Implications for AI Accountability: If interpretable eigenvectors genuinely capture the model’s decision-making features, they provide a basis for:

- Auditing model behavior for biases
- Explaining individual predictions
- Identifying potential failure modes

6.5 Communication with Original Authors

[Document any communication with original authors, or note that no communication occurred.]

7 Conclusion

[To be completed after experiments]

We reproduced both main experimental sections from Pearce et al. (2024): Section 4 (Vision), verifying that regularization induces low-rank interpretable structure in bilinear MLPs on MNIST; and Section 5 (Language), analyzing bilinear attention in GPT-2 and comparing with Sparse Autoencoders. Our ablation study [confirms/reveals] the relative importance of noise augmentation versus weight decay. The CP-decomposition extension demonstrates that structural low-rank constraints [can/cannot] achieve comparable interpretability to regularization-induced low-rank.

Limitations:

- The interpretability of eigenvectors is assessed qualitatively; future work could develop quantitative metrics.
- The language experiments focus on GPT-2; scaling to larger models remains future work.

Future Work:

- Apply bilinear decomposition to larger-scale vision models
- Develop automated methods for assessing eigenvector interpretability
- Explore other tensor decompositions (Tucker, TT) as alternatives to CP

Broader Impact Statement

This work contributes to the field of mechanistic interpretability, which aims to make neural networks more transparent and understandable. Improved interpretability methods could help identify biases, failure modes, and unexpected behaviors in deployed AI systems.

Author Contributions

Acknowledgments

References

- Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT Press, 2016.
- Tamara G Kolda and Brett W Bader. Tensor decompositions and applications. *SIAM review*, 51(3):455–500, 2009.
- Tim Pearce, Tuomas Oikarinen, and Tsui-Wei Weng. Bilinear mlps enable weight-based mechanistic interpretability. *arXiv preprint arXiv:2410.08417*, 2024.
- Olivier Roy and Martin Vetterli. The effective rank: A measure of effective dimensionality. In *European Signal Processing Conference (EUSIPCO)*, pp. 606–610, 2007.

A Additional Figures

B Hyperparameter Details

C CO2 Tracking Details