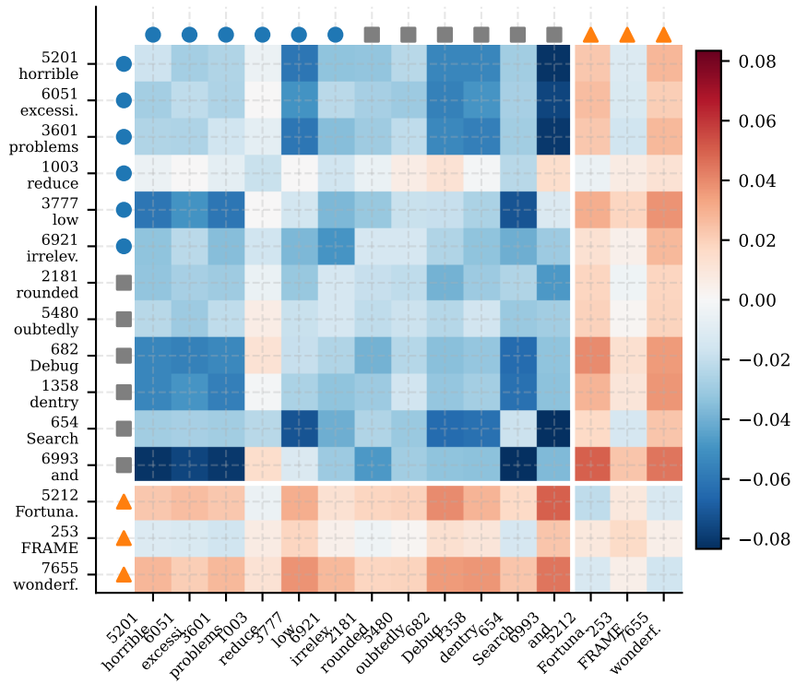
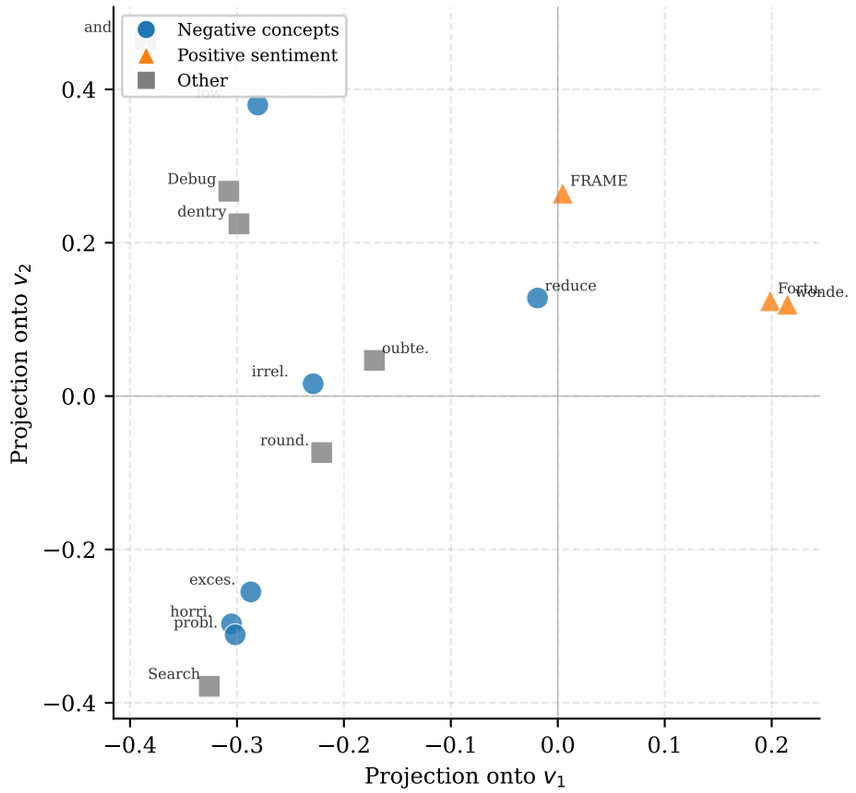
Sentiment-Related Circuit: Feature 751 (fw-medium, layer 7)
Per-token activations show which words activate each cluster

A) Interaction Submatrix (with semantic labels)



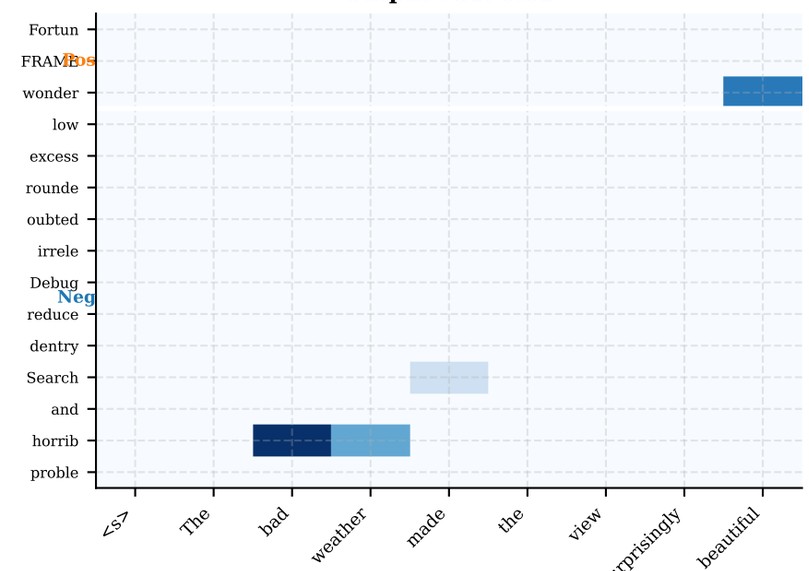
B) Feature Projections (all features labeled)



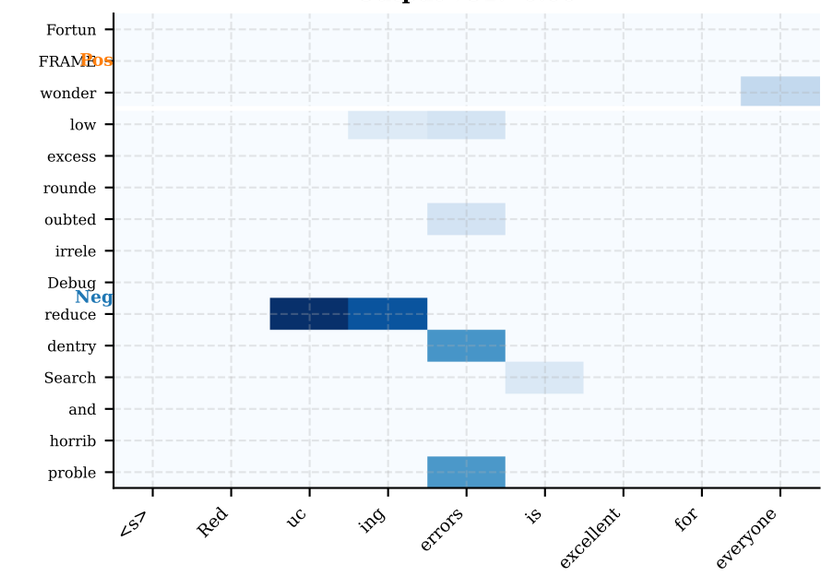
C) Feature Semantics (top 3 tokens per feature)

Feature	Cluster	Top Tokens
5201	Negative	'horrible', 'bad', 'terrible'
6051	Negative	'excessive', 'unfair', 'excess'
3601	Negative	'problems', 'illness', 'disease'
1003	Negative	'reduce', 'reducing', 'reduces'
3777	Negative	'low', 'few', 'lowest'
6921	Negative	'no', 'irrelevant', 'barely'
2181	Negative	'rounded', 'designers', 'of'
5480	Negative	'oubtedly', 'ailable', 'embr'
682	Negative	'`);', 'Debug', 'scription'
1358	Negative	'dentry', 'ishop', 'Nicole'
654	Negative	'V', 'U', 'Search'
6993	Negative	'HC', 'and', 'Department'
5212	Positive	'Fortunately', 'unately', 'provider'
253	Positive	'FRAME', 'down'
7655	Positive	'wonderful', 'spectacular', 'enjoyable'

Output 751: 0.62



Output 751: -0.00



Output 751: 0.21

